

# Neo4j Migration Guide: AuraDB Free → Self-Hosted Community Edition

## Project Information

Field	Value
Project	Weavink - NFC Business Card Platform
Author	Leo (CTO, Weavink)
Date	November 30, 2025
Server	Hetzner CX43 (8 vCPU, 16GB RAM, 160GB SSD)
Deployment Platform	Coolify
Neo4j Version	2025.10.1 Community Edition

## Table of Contents

1. [Executive Summary](#)
2. [Why Self-Host Neo4j?](#)
3. [Infrastructure Setup](#)
4. [Performance Benchmarks](#)
5. [The Migration Attempt That Failed](#)
6. [The Solution That Worked](#)
7. [Step-by-Step Migration Guide](#)
8. [Post-Migration Configuration](#)
9. [Troubleshooting Reference](#)
10. [Maintenance Commands](#)
11. [Scaling Guide](#)
12. [Lessons Learned](#)

## 1. Executive Summary

### What We Did

Migrated Neo4j graph database from **AuraDB Free** (Neo4j's managed cloud service) to a **self-hosted Neo4j**

**Community Edition** running in Docker on our Hetzner VPS via Coolify.

## Final Result

- **110 Contacts, 50 Tags, 190 Events, 12 Companies** successfully migrated
- **5-7ms query latency** (self-hosted) vs **30-70ms** (AuraDB) - **6-14x faster!**
- **€0/month** vs potential paid tier when exceeding free limits
- Full control over data and configuration

## Key Lesson

**AuraDB backups use Enterprise Edition format and CANNOT be restored to Community Edition.** You must export data using Cypher/APOC instead.

---

## 2. Why Self-Host Neo4j?

### AuraDB Free Tier Limitations

Aspect	AuraDB Free	Self-Hosted
<b>Nodes Limit</b>	200,000	Unlimited (disk-limited)
<b>Relationships</b>	400,000	Unlimited
<b>Cost</b>	€0 (with limits)	€0 (included in VPS)
<b>Latency</b>	~30-70ms	~5ms (localhost)
<b>Backups</b>	1 snapshot only	Configure yourself
<b>Data Location</b>	Google Cloud	Your server (GDPR)

### When to Migrate

- Approaching AuraDB limits (200K nodes)
- Need sub-10ms query latency
- Want full data sovereignty
- Running other services on same VPS anyway

---

## 3. Infrastructure Setup

### Server Specifications

Provider: Hetzner  
Model: CX43  
vCPU: 8

RAM: 16GB  
Storage: 160GB SSD  
Location: Falkenstein, Germany (EU)  
Cost: €8.99/month

## Neo4j Resource Allocation (Current Production Config)

Page Cache: 4GB (caches data for fast reads)  
Heap Initial: 2GB (query processing)  
Heap Max: 2GB (query processing)  
Total Reserved: ~6GB RAM

## Docker Compose Configuration (Coolify)

```
yaml
services:
  neo4j:
    image: 'neo4j:community'
    restart: unless-stopped
    environment:
      - NEO4J_AUTH=neo4j/YourSecurePassword123!
      - NEO4J_server_memory_pagecache_size=4G
      - NEO4J_server_memory_heap_initial_size=2G
      - NEO4J_server_memory_heap_max_size=2G
    volumes:
      - 'neo4j-data:/data'
      - 'neo4j-logs:/logs'
    healthcheck:
      test:
        - CMD
        - wget
        - '-q'
        - '--spider'
        - 'http://localhost:7474'
      interval: 30s
      timeout: 10s
      retries: 3
    volumes:
      neo4j-data: null
      neo4j-logs: null
```

## Ports (Internal Only)

- 7474: HTTP Browser (internal)

- **7687:** Bolt protocol (app connections)

**⚠️ Security:** Keep Neo4j internal to Docker network. No public exposure needed - your app connects via internal network.

## 4. Performance Benchmarks

### Test Environment

- **AuraDB:** Free tier, cloud-hosted
- **Self-Hosted:** Hetzner CX43, 4GB page cache, 2GB heap
- **Data:** 110 Contacts, 50 Tags, 190 Events, 12 Companies (~362 nodes)

### Results (Warm Cache)

Query	AuraDB	Self-Hosted	Improvement
<code>MATCH (c:Contact)-[:WORKS_AT]-&gt;(comp:Company)</code>	69-72ms	5ms	<b>14x faster</b>
<code>MATCH (c:Contact) RETURN count(c)</code>	22-23ms	3-4ms	<b>6x faster</b>
<code>MATCH (c:Contact)-[:SIMILAR_TO]-&gt;(other:Contact)</code>	31-33ms	5ms	<b>6x faster</b>

### Cold vs Warm Cache

State	Query Time	Notes
Cold (after restart)	~480ms	First query loads data into page cache
Warming	~12ms	Subsequent queries
Warm	<b>2-5ms</b>	Fully cached, production performance

### Understanding Neo4j Memory

Component	Location	Purpose
<b>Data on disk</b>	Docker volume	Permanent storage (~542MB)
<b>Page Cache</b>	RAM (4GB)	Caches disk data for fast reads
<b>Heap</b>	RAM (2GB)	Query processing and operations

The flow: `Disk → Page Cache (RAM) → Query Results`

Once your data is in the page cache, all queries run from RAM!

## 5. The Migration Attempt That Failed

### What We Tried First

## 1. Downloaded AuraDB Backup

- Went to AuraDB Console → Instance → Snapshots
- Downloaded `.backup` file (~100KB)
- File: `neo4j-2025-11-30T03-58-03-9077f1b7.backup`

## 2. Attempted Restore to Community Edition

```
bash

# Copied backup to server
scp ~/Downloads/neo4j-*.backup root@159.69.215.143:/root/

# Copied into container
docker cp /root/neo4j-*.backup neo4j-container:/var/lib/neo4j/backups/

# Fixed permissions
docker exec neo4j-container chown -R neo4j:neo4j /var/lib/neo4j/backups

# Attempted restore
docker exec -u neo4j neo4j-container neo4j-admin database load \
--from-path=/var/lib/neo4j/backups \
--overwrite-destination=true neo4j
```

## The Error

```
Files: 63/63, data: 100.0%
Done: 63 files, 1.924MiB processed in 0.099 seconds.
Failed to load database 'neo4j': Block format detected for database neo4j
but unavailable in this edition.
Load failed for databases: 'neo4j'
```

## Root Cause

**AuraDB (even the Free tier) runs on Neo4j Enterprise Edition internally.** When you export a backup, it uses the Enterprise "block format" which is **incompatible with Community Edition**.

This is NOT documented clearly by Neo4j. The backup downloads successfully, the restore runs to 100%, but fails at the final step.

## What This Corrupted

The failed restore attempt left the database in a broken state:

- The `neo4j` database became unavailable
- Only the `system` database remained

- Community Edition cannot create new databases (`CREATE DATABASE` is Enterprise-only)
- 

## 6. The Solution That Worked

### The Correct Approach: APOC Cypher Export

Instead of using the binary backup, we exported all data as **Cypher statements** using APOC (A Package Of Components), which is available in AuraDB.

### Why This Works

- Cypher is plain text, edition-agnostic
  - Creates `CREATE` statements for all nodes and relationships
  - Can be run on any Neo4j edition
  - Preserves all properties and relationships
- 

## 7. Step-by-Step Migration Guide

### Phase 1: Export from AuraDB

#### Step 1.1: Check Your Data Structure

Connect to AuraDB via Neo4j Browser and run:

```
cypher
-- List all node labels
CALL db.labels() YIELD label RETURN label

-- List all relationship types
CALL db.relationshipTypes() YIELD relationshipType RETURN relationshipType
```

Our results:

- Labels: Contact, Company, Tag, Event
- Relationships: WORKS\_AT, HAS\_TAG, SIMILAR\_TO, KNOWS, ATTENDS, MATCHED\_AT

#### Step 1.2: Export Using APOC

Run this query in AuraDB Browser:

```
cypher
```

```
CALL apoc.export.cypher.all(null, {format: "plain", stream: true})
YIELD cypherStatements
RETURN cypherStatements
```

### Step 1.3: Download as CSV

1. Click the **Export** button in Neo4j Browser
2. Choose **CSV** format
3. Save as `neo4j_export.csv`

The CSV contains all your data as executable Cypher statements.

---

## Phase 2: Deploy Neo4j Community Edition

### Step 2.1: Create Neo4j Service in Coolify

1. Go to **Coolify** → Your Project → + **Add Resource**
2. Select **Docker Compose** (NOT Dockerfile)
3. Paste the docker-compose configuration from Section 3
4. Click **Save** and **Deploy**

### Step 2.2: Verify Deployment

```
bash
# SSH to server
ssh root@159.69.215.143

# Check container is running
docker ps | grep neo4j

# Check logs
docker logs neo4j-container-name --tail 20
```

You should see:

```
INFO Bolt enabled on 0.0.0.0:7687.
INFO HTTP enabled on 0.0.0.0:7474.
INFO Started.
```

## Phase 3: Prepare Import File

### Step 3.1: Transfer CSV to Server

```
bash

# From your local machine
scp ~/Downloads/neo4j_export.csv root@159.69.215.143:/root/
```

### Step 3.2: Convert CSV to Cypher Script

The CSV has a header and quoted content. Clean it up:

```
bash

# SSH to server
ssh root@159.69.215.143

# Convert CSV to clean Cypher file
tail -n +2 /root/neo4j_export.csv | sed 's/^"/;s/"$/ | sed 's/""//g' > /root/neo4j_import.cypher

# Verify the file looks correct
head -5 /root/neo4j_import.cypher
```

You should see clean Cypher statements:

```
cypher

CREATE RANGE INDEX FOR (n:Contact) ON (n.userId);
CREATE CONSTRAINT company_name FOR (node:Company) REQUIRE (node.name, node.userId) IS UNIQUE;
...
```

---

## Phase 4: Import Data

### Step 4.1: Copy Import File to Container

```
bash

docker cp /root/neo4j_import.cypher neo4j-container-name:/var/lib/neo4j/
```

### Step 4.2: Run the Import

```
bash
```

```
docker exec -it neo4j-container-name cypher-shell \
-u neo4j \
-p 'YourSecurePassword123!' \
-f /var/lib/neo4j/neo4j_import.cypher
```

### Step 4.3: Verify Import

```
bash

# Count all nodes by type
docker exec -it neo4j-container-name cypher-shell \
-u neo4j -p 'YourSecurePassword123!' \
"MATCH (n) RETURN labels(n) AS type, count(*) AS count"

# Sample some data
docker exec -it neo4j-container-name cypher-shell \
-u neo4j -p 'YourSecurePassword123!' \
"MATCH (c:Contact) RETURN c.name, c.company LIMIT 5"
```

Expected output:

type	count
["Contact"]	110
["Tag"]	50
["Event"]	190
["Company"]	12

## 8. Post-Migration Configuration

### Update Weavink Environment Variables

In Coolify, update your Weavink app's environment variables:

#### Before (AuraDB):

```
env

NEO4J_URI=neo4j+s://xxxx.databases.neo4j.io
NEO4J_USERNAME=neo4j
NEO4J_PASSWORD=your-aura-password
```

## After (Self-hosted):

```
env  
  
NEO4J_URI=bolt://neo4j-ws4www84wss8scck0g0cow04:7687  
NEO4J_USERNAME=neo4j  
NEO4J_PASSWORD=YourSecurePassword123!
```

**Note:** Use `bolt://` not `neo4j+s://` for local connections. The container name is the hostname within Docker's network.

## Redeploy Weavink

After updating environment variables, redeploy the Weavink application in Coolify.

## Verify Memory Configuration

After any configuration change, verify settings:

```
bash  
  
docker exec neo4j-ws4www84wss8scck0g0cow04 cat /var/lib/neo4j/conf/neo4j.conf | grep memory
```

Expected output:

```
server.memory.pagecache.size=4G  
server.memory.heap.max_size=2G  
server.memory.heap.initial_size=2G
```

## 9. Troubleshooting Reference

### Problem: "Block format detected but unavailable in this edition"

**Cause:** Trying to restore Enterprise backup to Community Edition

**Solution:** Use APOC Cypher export instead of binary backup

---

### Problem: "Unable to get a routing table for database 'neo4j'"

**Cause:** Database corrupted or not created

**Solution:** Reset the data volume:

```
bash
```

```
# Stop container
docker stop neo4j-container-name

# Remove container (so volume can be deleted)
docker rm neo4j-container-name

# Find volume name
docker volume ls | grep neo4j

# Remove volume
docker volume rm volume-name-here

# Redeploy from Coolify
```

---

**Problem: "CREATE DATABASE is not supported in community edition"**

**Cause:** Community Edition only supports one database

**Solution:** Don't try to create databases. Use the default `(neo4j)` database.

---

**Problem: "Connection refused" when connecting to cypher-shell**

**Cause:** Neo4j still starting up

**Solution:** Wait 20-30 seconds after container start:

```
bash

sleep 30
docker logs neo4j-container-name --tail 10
# Should see "Started." before connecting
```

---

**Problem: Permission errors with neo4j user**

**Cause:** Running commands as root but Neo4j runs as `(neo4j)` user

**Solution:** Use `(-u neo4j)` flag or fix permissions:

```
bash
```

```
# Run as neo4j user
docker exec -u neo4j neo4j-container-name neo4j-admin ...

# Or fix permissions
docker exec neo4j-container-name chown -R neo4j:neo4j /var/lib/neo4j/
```

### Problem: Slow first query after restart

Cause: Cold cache - data needs to load from disk to page cache

Solution: This is normal. First query may take 400-500ms. Subsequent queries will be 2-5ms once cache is warm. Your app's regular queries will keep the cache warm.

## 10. Maintenance Commands

### Daily Operations

```
bash

# Check container status
docker ps | grep neo4j

# View recent logs
docker logs neo4j-ws4www84wss8scck0g0cow04 --tail 50

# Check database health
docker exec -it neo4j-ws4www84wss8scck0g0cow04 cypher-shell \
-u neo4j -p 'YourSecurePassword123!' \
"CALL dbms.components() YIELD name, versions RETURN name, versions"
```

### Check Disk Usage

```
bash

# Check data volume size
docker system df -v | grep neo4j

# Or check inside container
docker exec neo4j-ws4www84wss8scck0g0cow04 du -sh /data
```

### Backup (Manual)

```
bash
```

```

# Create a Cypher export (same method we used to migrate)
docker exec -it neo4j-ws4www84wss8scck0g0cow04 cypher-shell \
-u neo4j -p 'YourSecurePassword123!' \
"CALL apoc.export.cypher.all('/var/lib/neo4j/backup.cypher', {})"

# Copy backup out of container
docker cp neo4j-ws4www84wss8scck0g0cow04:/var/lib/neo4j/backup.cypher /root/backups/

```

**Note:** APOC may not be installed in Community Edition by default. You may need to add the APOC plugin to your Docker configuration.

## Query Statistics

```

bash

# Count all nodes
docker exec -it neo4j-ws4www84wss8scck0g0cow04 cypher-shell \
-u neo4j -p 'YourSecurePassword123!' \
"MATCH (n) RETURN count(n) AS totalNodes"

# Count all relationships
docker exec -it neo4j-ws4www84wss8scck0g0cow04 cypher-shell \
-u neo4j -p 'YourSecurePassword123!' \
"MATCH ()-[r]->() RETURN count(r) AS totalRelationships"

# Database size (approximate)
docker exec neo4j-ws4www84wss8scck0g0cow04 du -sh /data

```

## Restart Neo4j

```

bash

# Via Docker
docker restart neo4j-ws4www84wss8scck0g0cow04

# Via Coolify
# Go to Coolify UI → Project → Neo4j → Restart

```

# 11. Scaling Guide

## Current Resource Usage

Resource	Used	Allocated
Disk	~542MB	160GB

Resource	Used	Allocated
Page Cache	~542MB	4GB
Heap	Variable	2GB

## Capacity Planning

Users	Est. Nodes	Est. Disk	Fits in 4GB Cache?
1	~362	~542MB	✓ Yes
5-6	~2,000	~1-1.5GB	✓ Yes
50	~18,000	~5-8GB	⚠ Partially
200	~72,000	~20-30GB	✗ Need more RAM

## Memory Configuration Recommendations

User Count	Page Cache	Heap	Total RAM
1-50	4GB	2GB	~6GB
50-100	6GB	2GB	~8GB
100-200	8GB	3GB	~11GB

## How to Change Memory

Update docker-compose in Coolify:

```
yaml
environment:
  - NEO4J_server_memory_pagecache_size=6G # Increase for more data caching
  - NEO4J_server_memory_heap_initial_size=2G
  - NEO4J_server_memory_heap_max_size=2G
```

Then **Save** and **Redeploy**.

## When to Upgrade Server

If your data exceeds 10GB and you need fast queries, consider:

- Hetzner CX43 → CX53 (16GB → 32GB RAM, ~€18/month)
- Or optimize data model to reduce node/relationship count

## 12. Lessons Learned

### 1. AuraDB Backups ≠ Community Compatible

The binary (.backup) files from AuraDB use Enterprise Edition format. **Always use Cypher/APOC export for cross-edition migration.**

### 2. Test Before Production

We should have tested the restore on a local Docker instance before attempting on production server.

### 3. Community Edition Limitations

- Single database only (no CREATE DATABASE)
- No clustering
- Single neo4j user
- Some APOC procedures unavailable

### 4. Container Naming in Coolify

Coolify generates container names like neo4j-ws4www84wss8scck0g0cow04. Use docker ps | grep neo4j to find it.

### 5. Connection String Differences

- AuraDB: neo4j+s:// (encrypted)
- Self-hosted local: bolt:// (internal network, no encryption needed)

### 6. Failed Imports Corrupt Data

A failed import can leave the database in a broken state. Only solution is to delete the data volume and start fresh.

### 7. Cache Warmup is Normal

First query after restart may take 400-500ms. This is normal - Neo4j is loading data into the page cache. Regular app traffic keeps the cache warm.

### 8. Self-Hosted is Much Faster

With proper configuration, self-hosted Neo4j is **6-14x faster** than AuraDB due to:

- Zero network latency (localhost vs internet round-trip)
- Dedicated resources (no multi-tenant overhead)
- Data fits entirely in page cache

# Quick Reference Card

## Container Name

```
bash
```

```
docker ps | grep neo4j  
# Current: neo4j-ws4www84wss8scck0g0cow04
```

## Connect to Database

```
bash
```

```
docker exec -it neo4j-ws4www84wss8scck0g0cow04 cypher-shell -u neo4j -p 'YourSecurePassword123!'
```

## Connection String for App

```
bolt://neo4j-ws4www84wss8scck0g0cow04:7687
```

## Server Details

```
IP: 159.69.215.143
```

```
SSH: ssh root@159.69.215.143
```

## Check Memory Config

```
bash
```

```
docker exec neo4j-ws4www84wss8scck0g0cow04 cat /var/lib/neo4j/conf/neo4j.conf | grep memory
```

## Check Disk Usage

```
bash
```

```
docker system df -v | grep neo4j
```

## Useful Cypher Queries

```
cypher
```

```
-- Count all nodes
MATCH (n) RETURN labels(n), count(*)
```

```
-- List all indexes
SHOW INDEXES
```

```
-- List all constraints
SHOW CONSTRAINTS
```

```
-- Test query performance (run twice for warm cache)
MATCH (c:Contact)-[:SIMILAR_TO]->(other:Contact) RETURN c.name, other.name LIMIT 20;
```

## Document History

Date	Version	Changes
2025-11-30	1.0	Initial migration and documentation
2025-11-30	1.1	Updated memory config (4GB page cache, 2GB heap), added performance benchmarks, scaling guide

*Document created after successful migration from AuraDB Free to self-hosted Neo4j Community Edition on Hetzner VPS via Coolify.*